

The AB-Cloud as a Phase Resonator

Riemann ζ Zeros, GUE Universality, and Topological Vortex Matter: A Fully Verified Numerical Investigation

Isaev Iskhak Khamzatovich (ORCID: 0009-0003-7299-0701)

Version 22 · August 2026

Table of Contents

Abstract

This monograph has been written from scratch and summarizes a systematic numerical study of the object we call the **AB-cloud**: a lattice quantum system (a Hofstadter Hamiltonian with Aharonov–Bohm phases) in which topological vortex excitations generate an inhomogeneous phase texture, and the distribution of eigenvalues is compared against the statistics of the zeros of the Riemann zeta function. Every claim made here is accompanied by reproducible numbers from an automated 37-test verification suite (versions 18/19, August 2026). Every number carries its test ID; every test produces a full computation log; every suite run leaves on disk a directory with reports in four formats (Markdown, HTML, PDF, DOCX), figures in three raster/vector formats, and timestamped logs designed for independent verification.

The principal verified results are as follows. First, in the smooth gauge (the monumental vortex phase model), the mean ratio of adjacent level spacings of the AB-cloud spectrum on a torus reaches $\langle r \rangle = 0.585 \pm 0.026$ against the GUE target 0.5992 — a -2.4% deviation (Test 33) — and a scan over lattice sizes $L=10, 20, 30, 50$ shows monotone convergence $0.548 \rightarrow 0.592 \rightarrow 0.605 \rightarrow 0.595$ (Test 34). Second, a direct comparison of unfolded AB-cloud spacings with the first 5000 ζ zeros shows that the cloud’s pair correlation function is closer to GUE than to Poisson across the entire interval $s \in [0; 3)$, although the two-sample Kolmogorov–Smirnov test on so small an ensemble rejects exact distributional equality ($p=6.7 \cdot 10^{-4}$; Test 35) — we record this discrepancy honestly and explain it by ensemble finiteness. Third, the topological skeleton of the construction is confirmed at machine precision: the Dirac-string gauge yields flux $2\pi q$ through exactly one plaquette (Tests 15, 24); the Byers–Yang theorem holds to $3.5 \cdot 10^{-15}$ for integer charge and is violated by 0.34 for fractional charge (Test 25); the Chern number of the

lowest band is $C_1=2$ (Test 20); and Connes self-duality at $\alpha=1/2$ produces exactly 4 zero modes (Test 17). Finally, the low-energy dynamics exhibits a Dirac cone with linear gap scaling $E_{\min} \propto 1/L$, $R^2=0.9997$ (Test 19), a Dirac dip in the density of states with contrast $20\times$ at $\alpha=1/2$ (Test 30), and the Hatano-Nelson skin effect away from the Hermitian point $\sigma=1/2$ (Test 32).

Together these results form a coherent narrative: the AB-cloud is a phase resonator in which universal GUE statistics arise from time-reversal violation by complex hopping phases, vortices behave as relativistic particles, and the critical line $\sigma=1/2$ stands out as the self-dual point where the cloud's statistics come closest to those of the ζ zeros. We consistently separate the proven (machine precision), the reliably measured (calibrated statistical bands), and the still open (the large- T asymptotics), supplying every conclusion with an explicit criterion — the reader is never asked to take a number on faith.

The construction's reference-grade character is then stress-tested where it matters most: a dedicated large-scale scan over all vortex dials (Chapter 9, Test 34-ext) locates a broad optimal window — $N_v \approx 8-16$, $q \approx 0.7-1.0$, $W = 0.5-1.0$, strictly at the self-dual flux $\alpha = 0.5$ — in which the pooled KS distance to the unfolded spacings of 50,000 ζ zeros falls from $D = 0.1096$ (the reference configuration, WARN) to $D = 0.0318$ at the full 96×96 lattice: a factor of 3.4, closing $\approx 89\%$ of the distance to the finite-sample GUE floor (0.0220). We read the window as the vortex-density analogue of a universal physical hierarchy — a breeze needs fewer vortices than an ocean current — and as numerical support for the statistical arm of the Hilbert-Pólya program.

Keywords: Hilbert-Pólya conjecture; Riemann zeta function; random matrix theory; GUE; Aharonov-Bohm effect; Hofstadter Hamiltonian; AB-cloud; topological vortices; Dirac cone; TKNN; Hatano-Nelson skin effect.; vortex-density optimization; vortex matter

1. Introduction and Problem Statement

1.1 The Hilbert-Pólya Conjecture

In 1914 Riemann conjectured that all non-trivial zeros of the zeta function $\zeta(s)$ lie on the critical line $Re\,s=1/2$. Nearly half a century later, in 1958, David Hilbert (as recalled by Pólya) proposed the strategy now known as the **Hilbert-Pólya conjecture**: if one can find a self-adjoint operator $\hat{H}$ whose spectrum coincides with the zero ordinates $\{\gamma_n\}$, then the Riemann Hypothesis would follow immediately — Hermiticity forces all eigenvalues to be real, i.e. $Re\,\gamma_n=1/2$ for every n . The beauty of this program is that it translates analytic number theory into quantum mechanics: the zeta zeros become energy levels of some (as yet unknown) physical system.

The conjecture remains unproven, but it frames the present work: we construct and study a concrete class of quantum lattice systems whose spectral statistics approach universal random-matrix behavior. We emphasize the methodological position that distinguishes this monograph from much prior literature on the subject: we **do not claim a proof of the Riemann Hypothesis**. Our object is a precise numerical characterization of how closely finite lattice ensembles approach universal random-matrix statistics, and which parameters bring that convergence closer or push it away. All “closeness” criteria were fixed in advance, before the tests were run, and were never tuned along the way.

The historical context, briefly: in 1972 Montgomery formulated his pair-correlation conjecture for ζ zeros; Dyson, upon casually meeting Montgomery in the tea room of the Princeton Institute for Advanced Study, pointed out that Montgomery’s formula coincides with the pair correlation of eigenvalues of the GUE (Gaussian Unitary Ensemble) of random matrix theory. In the 1980s Odlyzko confirmed the coincidence numerically on an enormous dataset, and Berry in 1985 computed the finite-size corrections: a real (finite- T) sample of zeros must deviate from the pure GUE prediction in a predictable way — in particular, the pair correlation at the origin $R_2(0)$ receives a correction that decays as a power of T . These corrections are crucial for us: they explain why our Tests 4–6 and 11 see statistically significant deviations from GUE on samples with $T \lesssim 4 \cdot 10^4$, and why such deviations confirm rather than refute the picture.

1.2 The Aharonov–Bohm Effect and the AB-Cloud

The **Aharonov–Bohm (AB) effect** is the observation that a quantum particle responds to magnetic flux even where the field strength vanishes: a wave function encircling a solenoid carrying flux Φ acquires the phase $\varphi_{\text{AB}} = 2\pi\Phi/\Phi_0$, where $\Phi_0 = h/e$ is the flux quantum. On a lattice this is encoded by dressing the hopping amplitudes: $t_{ij} \rightarrow t_{ij} e^{i\varphi_{ij}}$. The **Hofstadter Hamiltonian** is the canonical instance: electrons on a square lattice in a uniform magnetic field with flux $\alpha = \Phi/\Phi_0$ per plaquette; its spectrum as a function of α is the famous Hofstadter butterfly (reproduced as Fig. 19 of our figure set).; Fig. 20 — the KS landscape $D(Nv, q)$ at $L = 24$ (Stage A); Fig. 21 — the size ladder $L = 24 \rightarrow 64$ with per- L GUE floors (Stages B–C)

By the **AB-cloud** we mean the generalization in which pointlike **topological vortices** are placed on the lattice — centers around which the hopping phase winds by $2\pi q$ (q is the vortex charge, which may be fractional). The phase field of such a system is inhomogeneous: near each vortex the phase curls, and the combined picture resembles a cloud — hence the name. Mathematically the Hamiltonian is

$$H_{ij} = -\exp(i\varphi_{ij}), \quad \varphi_{ij} = 2\pi\alpha\sum_{k=1}^{N_v} q_k [\arg(\mathbf{r}_i - \mathbf{r}_k) - \arg(\mathbf{r}_j - \mathbf{r}_k)] + \varphi^{\text{str}}_{ij},$$

where j is the lattice column index (Landau gauge), N_v the number of vortices, and $\varphi^{\text{str}}_{ij}$ the explicit **Dirac string**: an addition of $-2\pi q$ attached to every bond crossing the half-line (string) emanating from the vortex. The Dirac string guarantees that the total magnetic flux around the vortex equals exactly $2\pi q$ to machine precision — verified by direct computation (Tests 15, 24).

Beyond the geometric (phase) part, the model carries **on-site disorder** W : additive node noise modeling impurities in a real crystal. Model parameters: lattice size $L \times L$, flux α , charge set $\{q_k\}$, vortex number and positions N_v , disorder strength W , and boundary conditions (open or torus). A key physical switch is the **vortex phase model**: (a) the Dirac-string gauge, in which integer-charge vortices are spectrally invisible (an exact Byers–Yang theorem, Test 25), and (b) the smooth monumental gauge (an atan model), in which phases are spread smoothly and are complex for any q — this is what makes GUE statistics reachable at physically sensible parameters. Both gauges are implemented in the code and used as appropriate: flux theorems are verified in the first, spectral statistics in the second.

1.3 Connection to Random Matrix Theory

Random matrix theory (RMT) studies eigenvalue distributions of large matrices with random entries. Three ensembles matter for us, distinguished by time-reversal symmetry: **GOE** (real symmetric, T-symmetric), **GUE** (complex Hermitian, T-violating), and **GSE** (quaternionic self-dual), plus the limiting **Poisson** case (uncorrelated levels; integrable systems). The diagnostics are defined in Section 2.2; here we note the central scale of the mean adjacent-gap ratio: $\langle r \rangle \approx 0.386$ (Poisson), 0.531 (GOE), 0.600 (GUE). Dyson’s threefold way predicts that T-violation by complex phases necessarily moves a system into the GUE class — and the AB-cloud in the smooth gauge is a textbook illustration: at $\alpha=1/2$ the matrix is complex Hermitian with $\tau(H) = \|H - H^i\| / \|H\| = 0.225 \neq 0$ (Test 26), and the statistics indeed approach GUE, whereas in the real Dirac-string gauge $\tau=0$ and the spectrum is locked into the GOE class (Tests 15, 16).

1.4 Goals, Tasks, and Verification Organization

The goal of this work is a complete, independently checkable picture of AB-cloud spectral physics. The tasks are concretized as 37 automated tests grouped into six blocks: (1) convergence of the $b(N)$ correction and the reviewers’ objections (Tests 1–14); (2) construction and topological theorems (15, 17, 20, 24–27); (3) GUE classification of the cloud (16, 26, 29, 33, 34); (4) direct comparison with ζ zeros (28, 35, 36); (5) Dirac dynamics

and non-Hermiticity (19, 30–32, 37); (6) analytical interpretations (Tests 21–23 and Section 8). Each test returns a boolean verdict against a pre-registered criterion; verification is two-pass (first pass — the primary lattice, second — an enlarged one), and the “quick check” runs all 37 tests on $16 \times 16 \rightarrow 32 \times 32$ lattices with a ζ sample truncated to 5000 zeros in minutes.

We explicitly distinguish the **two verdict systems** so the reader never confuses them. Formulaic tests (identities: fluxes, self-duality, Byers–Yang) are checked with machine tolerances $10^{-12} \dots 10^{-16}$ and yield a binary PASS/FAIL. Statistical tests (GUE agreement) are calibrated against confidence bands established by simulation: e.g., PASS on $\langle r \rangle$ requires the band 0.5992 ± 0.022 ; these admit a third verdict, WARN — “a deviation was measured but is explained by finite size”. This dual system removes the cardinal sin of numerical work — tuning the criteria to the result.

2. Mathematical Background

2.1 The Riemann Zeta Function and Its Zeros

The **Riemann zeta function** $\zeta(s) = \sum_{n \geq 1} n^{-s}$ extends analytically to the whole complex plane and has non-trivial zeros in the critical strip $0 < \text{Re } s < 1$. Numerically we use the first 50,000 zeros (Odlyzko’s dataset), whose ordinates fill $T \in [14.13; 40433.7]$. The density of zeros at height T is given by the **Riemann-von Mangoldt formula**:

$$N(T) = \frac{T}{2\pi} \ln \frac{T}{2\pi} - \frac{T}{2\pi} + \frac{7}{8} + O(\ln T),$$

which grows roughly as $(T/2\pi) \ln T$. Dividing out the local density — the **unfolding** procedure — makes the normalized spacings $\delta_n = \gamma_{n+1}^{(u)} - \gamma_n^{(u)}$ have unit mean by construction, so all correlation information is carried by their distribution. In the code, unfolding uses a symmetric sliding window over the smooth $N(T)$; this standard step is also where artifacts most often creep in, which is why Test 10 checks the stability of all final quantities under subsampling (maximal relative deviation 8.5% against a 10% threshold — PASS).

We use zero statistics at three heights: the full 50,000-zero sample (Tests 4, 6–10), “high-temperature” subsamples with $T > T_{\min}$ (Test 5), and the first 1000/5000 zeros for two-pass comparisons (Tests 28, 35). The asymptotic nature of GUE statistics means agreement improves with T ; our data demonstrate this directly (Section 3.2).

2.2 Random-Matrix Ensembles and Diagnostic Statistics

All diagnostics are defined for unfolded (dimensionless) levels. We use five quantities, each sensitive to a different aspect of correlations.

Mean adjacent-gap ratio $\langle r \rangle$. For every triple of consecutive levels E_{n-1}, E_n, E_{n+1} set $r_n = \min(\delta_n, \delta_{n+1}) / \max(\delta_n, \delta_{n+1})$; the average is a quantity that requires no external unfolding (ratios are insensitive to local density). Reference values (large matrices): Poisson 0.3863, GOE 0.5307, GUE 0.5996. This is the primary diagnostic for finite lattice ensembles since it does not accumulate unfolding error.

Kolmogorov-Smirnov (KS) test. Compares the empirical CDF of spacings with the GUE Wigner prediction: the statistic D is the maximum absolute difference of the curves, and at n spacings the 5% critical value is $1.36/\sqrt{n}$. For $n=49999$ the critical $D=0.0061$ — even microscopic deviations “reject” GUE; hence we always report D together with n and the critical threshold (Test 4), and separately track the trend $D(T_{min})$ (Test 5).

Number variance $\Sigma^2(L)$ — the variance of the level count in a window of length L (in units of mean density). Poisson gives $\Sigma^2=L$; GUE gives $\Sigma^2 \approx (1/\pi^2) \ln(2\pi L) + \dots$ — logarithmic rather than linear growth. Spectral rigidity $\Delta_3(L)$ is the integral version: the mean-square deviation of the best straight line fitted to the spectral staircase in a window L ; Poisson $L/15$, GUE $\approx (1/\pi^2) \ln L$. Both are computed on a common grid of nine L values for the data and for a simulated GUE reference (Tests 13, 14).

Pair correlation $R_2(s)$ — the probability density of finding a level at distance s from a given one (the Montgomery-pair-correlation statistic): for GUE, $R_2(s) = 1 - (\sin \pi s / (\pi s))^2$, with the correlation hole $R_2(0)=0$ expressing level repulsion; for Poisson, $R_2 \equiv 1$. The spectral form factor $K(t)$ — the Fourier transform of R_2 , “ramp + plateau” $K(t) = \min(t, 1)$ for GUE — is the most integral of our diagnostics (Tests 35, 36).

Anderson-Darling A^2 — a KS analogue weighted toward the distribution tails; we estimate the p-value by Monte Carlo over 2000 simulations (Test 11). Finally, **f_GUE** is a summary “GUE share” measure: $f_{\text{GUE}} = (\Sigma^2_{\text{P}} - \Sigma^2_{\text{data}}) / (\Sigma^2_{\text{P}} - \Sigma^2_{\text{GUE}})$ evaluated at $L=2$; $f_{\text{GUE}}=1$ is perfect GUE, 0 is Poisson (Test 29).

2.3 The Hofstadter Hamiltonian with Vortices: Exact Definition

Collecting the full definition. Lattice $L_x \times L_y$; the wave function lives on sites; matrix elements $H_{ij} = -\exp(i\phi_{ij})$ for nearest neighbors (including periodic wraps if the torus is chosen). The Landau gauge assigns column j the phase

$2\pi\alpha j$. A vortex of charge q_k at $r_k=(x_k, y_k)$ (in cell units) adds to the bond phase difference ($i \rightarrow j$) the quantity $q_k[\arg(r_i - r_k) - \arg(r_j - r_k)]$, and the Dirac string adds the discrete $-2\pi q_k$ on bonds crossing the string; in the smooth gauge the string is absent, with the atan phase distributed over vertical bonds. The key property: **the sum of phases around any plaquette equals $2\pi\alpha + 2\pi q_k \delta_k$** (plus $2\pi \times \text{integer}$), where δ_k indicates that the plaquette contains the vortex. This was verified for all 224 interior plaquettes of a 16×16 lattice: $\max |flux|$ on empty plaquettes $1.5 \cdot 10^{-14}$ rad (Test 24).

On-site disorder W is realized as node potentials $\varepsilon_i \in [-W/2; W/2]$ with a fixed generator seed (reproducibility: same seed — same disorder). Importantly, for the topological identities (Byers–Yang) the disorder is switched off: any spectral dependence on q beyond pure gauge would be an error, and Test 25 confirms this to $3.5 \cdot 10^{-15}$.

Boundary conditions are a separate physical choice. Open boundaries produce the butterfly’s edge states, which admix a Poisson component and depress $\langle r \rangle$ (Test 33: 0.442 on open boundaries vs. 0.585 on the torus at otherwise identical parameters). The torus requires the gauge-compatibility condition $q(2 - N_x) \in \mathbb{Z}$ — periodicity of the phase under a wrap; for integer q it holds automatically, for fractional $q=0.3$ on a 16×16 torus it fails, and the system honestly falls back to open boundaries (auto-detection with a diagnostic message in the code).

2.4 Verification Methodology and Reproducibility

Every suite run creates `results/run_<timestamp>/` with a subdirectory per test: inside are the full computation log (every quantity timestamped, with deltas against the previous entry), the captured console output, reports in Markdown, HTML, vector PDF and DOCX, figures in PNG (600 dpi), SVG and PDF, and a frame-by-frame GIF animation of the first figure. The logs are designed so that an outside researcher can re-derive every number of this monograph without access to our code: formula, inputs, output, timing. The reproduction commands are `julia ab_cloud_v19.jl --test <N> --no-two-pass` for a single test, `--test all` for the full suite, `--quick` for the fast check ($16 \times 16 \rightarrow 32 \times 32$, 5000 zeros, both passes). The report engine is isolated from the computational core: failure of any format cannot affect a test verdict (guaranteed architecturally in v19; see Appendix A).

3. Block 1: The $b(N)$ Correction and ζ -Zero Statistics (Tests 1–14)

3.1 What $b(N)$ Is and Why Its Convergence Rate Matters

The correction $b(N)$ arises in the analytic treatment of zero-pair correlations: it is a fitting quantity connecting the finite-sample value of a diagnostic to

its asymptote. Earlier versions of the program matched it to a theoretical expectation $\sim N^{-1/2}$; reviewers correctly noted that this expectation was never derived. We accept the objection: since suite version 6 the exponent of $b(N)$ is a **purely empirical** quantity with no theoretical template, and all Block-1 tests check the internal consistency of measurements rather than agreement with a non-existent theorem. Tests 3, 7, and 9 provide three independent measurements of the same exponent — their agreement is the object of verification.

3.2 Three Independent Measurements of the Exponent




Fig. 1. *Convergence of the $b(N)$ correction: data and two competing fits (power law and $1/\ln N$). The critics' $N^{-1/2}$ expectation is excluded by the data.*

Log-log regression (Test 3) over the grid $N \in [50; 25600]$ gives $b(N) \approx 7.03 N^{-0.1685}$, $R^2=0.9895$, slope standard error 0.0043. The **extended regression up to** $N=32000$ (Test 7) gives slope -0.1504 with a 95% confidence interval $[-0.1594; -0.1414]$, $R^2=0.9954$ — consistent. The **bootstrap** (Test 9, 1000 resamples with replacement) gives a median slope of -0.1746 and a 95% interval $[-0.1895; -0.1552]$. A detail that usually escapes notice: the **alternative** parametrization $b(N) \approx -0.40 + 16.84/\ln N$ describes the data better than the power law ($R^2=0.9999$ vs. 0.9954 in Test 7). This is a serious argument that the true asymptotics is logarithmic, not power-law; we plot both fits (Figs. 1 and 3) and do not adjudicate between them, since on the available range of N they cannot be definitively separated. Residual analysis (Test 8) shows serial structure (3 runs against an expectation of 5.8 for randomness; lag-1 autocorrelation 0.68) — the data “bend” around the

power-law fit exactly as a logarithmic form prescribes, consistent with the two-fit picture.

The practical conclusion of the block: on samples $N \leq 5 \cdot 10^4$ the correction $b(N)$ remains of order unity (from 2.39 at $N=500$ down to 1.28 at $N=32000$) — finite-sample effects in ζ statistics are large, and any “GUE or not” test that ignores $b(N)$ is simply invalid at these heights.

3.3 Objection 2: χ^2 and A^2 Reject GUE — and That Is Expected




Fig. 2. Histogram of 49,999 ζ -zero spacings against the Wigner surmises (GUE, GOE) and the Poisson exponential: a typically “matrix-like” shape with level repulsion.




Fig. 3. The KS statistic D against the lower cutoff $T_{\min}$: a slow drift toward the critical line — the numerical expression of GUE asymptoticity.

Tests 4, 5, 6, and 11 state the reviewers' objection in its strongest form and test it honestly. On the full sample of 49,999 spacings, the KS statistic against GUE is $D=0.0881$ against a critical value of 0.0061 — a fourteenfold excess, H_0 (GUE) rejected with $p \approx 10^{-300}$. The chi-square test on histograms with 150/300/600 bins gives $\chi^2=6170/6276/6384$ at 149/295/550 degrees of freedom — rejection at all scales. Anderson-Darling on 5,000 spacings against the analytic GUE CDF gives $A^2=117.7$ against a Monte-Carlo 95th percentile of 16.6 ($p \approx 0$ over 2000 simulations).

It would be a mistake to read this as a refutation of the GUE hypothesis for the zeros — and here is why. GUE statistics is **asymptotic**: it sets in at heights $T \gg 10^6$ where millions of zeros are available. Berry (1985) computed the leading finite-size corrections, whose sign and magnitude match the observations: at our heights $T \leq 4 \cdot 10^4$ a value $D \sim 0.08$ is exactly what should be expected. Direct confirmation comes from Test 5 (the high- T cut: $D=0.0878 \rightarrow 0.0866$ as $T_{\min}: 1000 \rightarrow 10000$ — a slow but systematic decrease) and, most importantly, Test 12: the **two-sample** KS between real spacings and the spacings of a simulated GUE matrix 2000×2000 gives $D=0.0406$, $p=0.139$ — the hypothesis “zero spacings are distributed like the spacings of a same-size GUE matrix” is **not rejected**. In other words: the zeros are indistinguishable from a finite-size GUE matrix but distinguishable from the abstract infinite GUE limit — which is precisely what finite-dimensional GUE-likeness means.

3.4 Spectral Rigidity: The Zeros Are More Rigid Than Simulated GUE




Fig. 4. The two fits of the $b(N)$ decay in log coordinates: power law and $1/\ln N$; the systematic bend of residuals points to a logarithmic asymptote.




Fig. 5. Bootstrap distribution of the slope (1,000 replicates): the 95% CI is $[-0.190; -0.155]$; the value -0.5 lies far outside.




Fig. 6. Number variance $\Sigma^2(L)$: the data (flat) against logarithmic GUE and linear Poisson — closer to GUE in 9/9 points.




Fig. 7. Spectral rigidity $\Delta_3(L)$: super-rigidity of the data relative to even the simulated GUE at large L .

The most informative results of the block come from $\Sigma^2(L)$ and $\Delta_3(L)$ (Tests 13, 14). On a nine-point grid $L \in [2; 1400)$, the zero data give $\Sigma^2 \approx 0.33 - 0.41$ —

essentially constant — while the Poisson prediction grows linearly to 1400 and the simulated GUE reference grows logarithmically to 0.61–0.72. The data are closer to GUE than to Poisson in all 9/9 points; but they are **stiffer** (smaller) than even the simulated GUE at large L — the relative excess of the data's Δ_3 over the GUE reference falls from 0.99 at $L=2$ to 0.48 at $L=1400$. Such super-rigidity is a known feature of the actual ζ zeros (spectral rigidity of arithmetic systems exceeds the matrix one), and artifacts of unfolding would inflate Δ_3 and Σ^2 rather than deflate them — so their reproduction is an independent certificate of unfolding quality.

4. Block 2: AB-Cloud Construction and Topological Theorems (Tests 15, 17, 20, 24–27)

4.1 Correctness of the Construction: Hermiticity and Fluxes

Before speaking of statistics, the construction must be exact. Test 15 checks on a 16×16 lattice with two vortices $q=\pm 1$: Hermiticity $\|H - H^\dagger\| / \|H\| = 0$ (machine zero); the symmetry class — $\tau(H) = \|H - H^c\| / \|H\| = 0$, i.e. in the Dirac gauge the matrix is real (the GOE class, as Dyson's trichotomy predicts); the flux through each vortex plaquette and a control empty plaquette to 10^{-14} rad. Test 24 repeats the flux check for a single vortex over all 225 plaquettes: 1/1 vortex plaquettes OK, 224/224 empty OK, $\max |flux|$ on empty plaquettes $1.47 \cdot 10^{-14}$. These two checks are the foundation: any error of the phase construction would corrupt the fluxes, yet the fluxes are perfect.

4.2 Byers–Yang: Integer Vortices Are Invisible, Fractional Ones Are Physical




Fig. 8. The Byers–Yang test: integer charge is spectrally invisible ($3.5 \cdot 10^{-15}$), fractional charge is physical (0.34); log scale.

The Byers–Yang theorem (1961) in its lattice formulation states: the spectrum is periodic under the addition of an integer flux quantum per plaquette — an integer-charge vortex can be “gauged away” by a unitary site transformation changing only phases. Test 25 measures this directly: two random configurations of four neutral vortices $q=\pm 1$ on 16×16 with open boundaries at $\alpha=0$ give a maximal spectral shift of $3.49 \cdot 10^{-15}$ relative to the clean lattice — integer vortices are spectrally invisible, the identity is exact. A single fractional configuration $q=0.3$ gives a shift of 0.342 — a fractional vortex carries physical flux $4\pi \cdot 0.3 = 1.2\pi \neq 0 \pmod{2\pi}$ and genuinely perturbs the spectrum. The consequence governing all subsequent physics: in the Dirac gauge only **fractional** charges or a background $\alpha \neq 0$ can generate GUE statistics; in the smooth monumental gauge the phases are complex for any q , and integer charges work (Section 5). Both branches are verified numerically, and this fork explains much of the historical disagreement between suite versions.

4.3 Connes Self-Duality and All Chiral Symmetry

At $\alpha=1/2$ the Hofstadter Hamiltonian exhibits **Connes self-duality**: the replacement $\alpha \rightarrow 1/\alpha$ is equivalent to a Fourier transform over sites, and $\alpha=1/2$ is a fixed point. On a torus, Test 17 measures: the number of zero modes at $\alpha=1/2$ equals exactly 4 (for even L this is twice the number of fixed-point edges — the analytic value), and the spectral $E \rightarrow -E$ symmetry

has defect $1.4 \cdot 10^{-15}$. **Sublattice chiral symmetry** (Test 27): on the bipartite square lattice the operator Γ (the checkerboard sign) satisfies $\Gamma H \Gamma = -H$ for any α in the absence of vortices; the measured defect at $W=0$ is 0 (exact), at $W=4$ it is $5.9 \cdot 10^{-3}$ (small but non-zero, since on-site disorder breaks bipartiteness), and with two vortices without disorder it is again exactly 0: phase vortices preserve bipartiteness. An additional analytic fact from Test 17: at $\alpha=1/3$ the spectrum is gapped at $E=0$ (zero modes: 0), and the isospectrality $\alpha \leftrightarrow 1/\alpha$ on a finite torus is violated (defect 0.30) — it is exact only in the Bloch-band sense; this is documented honestly, since the finite-size violation does not contradict the infinite-volume theorem.

4.4 The Chern Number: $C_1 = 2$

Test 20 computes the first-band **Chern number** by the Fukui-Hatsugai-Suzuki method on a 6×6 torus at $\alpha=1/2$ (no vortices): an 8×8 grid of twisted boundary conditions, the Berry section of the lowest band gives $C_1=2$. The monograph predicted $|C_1|=1$ per sub-band; at $\alpha=1/2$ the Hofstadter bands pair up, and the “lowest band” contains two sub-bands with $C_1=\pm 1$ contributions of equal sign — the total $|C_1|=2$ agrees with the prediction once the degeneracy is properly counted. The essential point stands: $C_1 \neq 0$ — the phase realizes the quantum Hall effect, the topology is non-trivial, and protected edge states must exist (on open boundaries they appear as the butterfly’s “legs”, Fig. 19).

4.5 The Torus and Time-Reversal Violation

Test 26 combines flux and statistical checks on a periodic lattice: torus 16×16 , a single vortex $q=1$ under the gauge-compatible condition $q(2 - N_x) = -14 \in \mathbb{Z}$, a periodic Dirac string along $+x$. Fluxes: the vortex plaquette and all 255 empty ones — OK; Hermiticity — OK; in the smooth gauge used for statistics, $\tau(H)=0.2252$ — T-symmetry is broken (the GUE class), and $\langle r \rangle = 0.5857$, inside the calibrated band around GUE 0.5992. This test is a microcosm of the whole architecture: exact fluxes + broken T + GUE statistics in a single configuration.

5. Block 3: GUE Classification of the Cloud (Tests 16, 26, 29, 33, 34)

5.1 The Role of the Gauge: Why the Same Model Yields GOE and GUE

Test 16 (the Montgomery test) in the Dirac gauge at $N_v=2$, $q=\pm 1$, $\alpha=1/2$ gives $\langle r \rangle = 0.5133 \pm 0.0353$ — between GOE (0.531) and GUE (0.600), with a Monte-Carlo p-value of 0.00 against a GUE ensemble of 30 matrices 256×256 . The reason is clear from 4.2: integer charges in this gauge are gauge-invisible, the matrix is nearly real, and the GOE “ceiling” cannot be overcome by any disorder — this is matrix-level, not noise-level T-symmetry.

Test 16 is repeated in the smooth gauge inside Test 33: at the same vortex positions $\langle r \rangle$ rises to 0.585. The moral for reproducibility: **the phase gauge is part of the physical specification**; comparing suite versions without stating the gauge is meaningless, and every report header now states it.

5.2 The Multi-Realization Bootstrap (Test 33)




Fig. 9. Test 33: three $\langle r \rangle$ realizations for $q=1$ (torus) and $q=0.3$ (open boundaries) against the GUE and Poisson lines.

The principal statistical result of this work. Configuration: torus 16×16 , two vortices $q=\pm 1$, $\alpha=1/2$, disorder $W=4$, smooth gauge, three independent realizations (random vortex positions, distinct disorder seeds). Results: $\langle r \rangle = 0.5639, 0.5811, 0.6094$; combined $\langle r \rangle = 0.5848 \pm 0.0260$, 95% CI $[0.5588; 0.6108]$, deviation from the GUE target 0.5992 equal to -0.0144 (-2.4%) — inside the calibrated band, PASS. A control series with fractional charge $q=\pm 0.3$ on open boundaries: $0.4025, 0.4967, 0.4260 \Rightarrow 0.4417 \pm 0.0555$ — a substantial depression explained by edge states (the Poisson component) and the absence of torus gauge compatibility. The gap between the two series (0.585 vs. 0.442) is the most vivid demonstration of how boundary conditions affect finite-lattice statistics (Fig. 11).

5.3 Scaling with Lattice Size (Test 34)




Fig. 10. Test 34: convergence of $\langle r \rangle$ with lattice size; both branches settle onto the GUE plateau ≈ 0.60 .

The scaling $\langle r \rangle(L)$ for $L=10, 20, 30, 50$ (two realizations per point, disorder $W=4$): the fractional branch $q=0.3$ (open boundaries) — $0.5482 \rightarrow 0.5921 \rightarrow 0.6048 \rightarrow 0.5953$; the integer branch $q=1$ (torus) — $0.5739 \rightarrow 0.5681 \rightarrow 0.6042 \rightarrow 0.5864$. Both branches reach the plateau 0.59 – 0.605 at $L \geq 30$, within two standard errors of GUE 0.5996 ; the fractional branch is monotone and its plateau matches the monograph’s prediction ($0.487 \rightarrow 0.562 \rightarrow 0.595 \rightarrow 0.594$) in convergence shape — our $L=10$ value is higher (edge effects at small L are weaker in our vortex configuration), and beyond that the trends coincide. This resolves the “decay paradox” of earlier versions: $\langle r \rangle$ does not degrade with lattice growth but settles onto the GUE plateau; the depression seen in v13 was an artifact of small sizes and the real gauge.

5.4 The Summary Measure f_{GUE} (Test 29)

Test 29 compresses the statistics onto one scale: at 16×16 , $N_v=2$, $q=\pm 1$, $\alpha=1/2$, $W=4$, the value $\Sigma^2(L=2)=0.8377$ yields $f_{\text{GUE}} = 0.648$ against a PASS threshold of 0.90 — WARN. This does not contradict Test 33: Σ^2 at small L is sensitive to the stiff spectral core and the edges, where the mixture is not yet universal; the 0.90 threshold was calibrated for ensembles $\geq 100 \times 100$. In Block 4 we will see the same picture in terms of $K(t)$: small ensembles show the right *trend* but fail hard thresholds calibrated on much

larger systems. We regard this spread of indicators not as a failure but as a measurable characteristic of finite size, recorded in the summary table.

6. Block 4: Direct Comparison with the ζ Zeros (Tests 28, 35, 36)

6.1 The Two-Pass Zero-Sample Comparison (Test 28)




Fig. 11. The Berry correction $R_2(0; T)$: the finite-sample cutoff and the slow approach to the -1 asymptote; verticals mark the two pass heights (1000/5000 zeros).

Test 28 — the Berry correction to $R_2(0)$ — runs in v18/v19 in two passes: the first 1000 zeros ($T_{max} \approx 1419$) and the first 5000 zeros ($T_{max} \approx 5448$). The prime sum $\sum_p \sum_{k \leq 5} p^{-k/2} / (k \ln p) = 4.2079$ gives the corrected theoretical value $R_2(0) = -2.3394$; the expected finite-sample addition $\sim T^{-1/2}$ decreases from 0.0019 ($T = 1419$) to 0.0009 ($T = 5448$). The test verifies the exact reproduction of the formula in both passes (PASS, machine precision) and demonstrates the scale of the Berry effect: even the corrected prediction for $R_2(0)$ at reachable heights is far from the textbook -1.007 asymptote. For the suite, the practical meaning is that the two passes set the scale on which any “agreement” of small zero samples must be interpreted.

6.2 The Montgomery Test: AB-Cloud vs ζ -5000 (Test 35)




Fig. 12. Pair correlation $R_2(s)$: the AB-cloud against the ζ zeros, GUE and Poisson — the correlation hole is reproduced by both branches.

The direct head-on comparison: unfolded AB-cloud spacings (2 realizations, 304 gaps) against the unfolded spacings of the first 5000 ζ zeros (4999 gaps, symmetric sliding-window unfolding). Results: two-sample KS $D=0.1173$, $p=6.7 \cdot 10^{-4}$ — the hypothesis of full distributional equality is rejected; the mean absolute difference of pair correlations $|\langle R_2^{\text{AB}}(s) - R_2^{\zeta}(s) \rangle| = 0.147$ against a threshold of 0.10. But the same $R_2(s)$ table (29 values of s from 0.05 to 2.85; Fig. 8) shows the essential point: the correlation hole exists for both — $R_2^{\text{AB}}(0.05) = 0.00$, $R_2^{\zeta}(0.05) = 0.005$ against the Poisson value of one; both curves run parallel to the GUE prediction, and both pass the “closer to GUE than to Poisson” check ($d_{\text{GUE}} = 0.140 < d_{\text{Pois}} = 0.227$). The honest summary: **the AB-cloud reproduces the qualitative and semi-quantitative correlation structure of the zeros (the hole, the kink near $s \approx 1$, the plateau), but does not match them pointwise** — 304 gaps are insufficient for a stable R_2 or a KS-level match. We record the distance between “the same universality” and “the same spectrum” as a number, not an adjective.

6.3 The Form Factor $K(t)$ (Test 36)




Fig. 13. Spectral form factor $K(t)$: the GUE ramp+plateau, the AB-cloud data and the ζ zeros; small ensembles show the right trend with outliers.

The spectral form factor is the most demanding diagnostic: a stable estimate requires hundreds of independent gaps and a precise unfolding. Our 306 eigenvalues give an RMS deviation from the GUE ramp of 0.93 (PASS threshold 0.30) and a correlation with the theoretical curve of 0.018 (threshold 0.50) — WARN on both metrics; the shape of the curve (Fig. 9) shows the correct rise toward $t \approx 1$ followed by a falloff, with outliers typical of small ensembles. The ζ -zero curve on the same plot (from 5000 zeros) has RMS against GUE of 0.86 — that is, even the reference zeros at these scales fail the hard threshold: the diagnostic is informative but calibrated for ensembles orders of magnitude larger. We include it in the monograph as a demonstration of applicability limits and as a task for future large runs.

7. Block 5: Dirac Dynamics and Non-Hermiticity (Tests 19, 30–32, 37)

7.1 The Dirac Cone (Tests 19, 31)




Fig. 14. The Dirac cone: $E_{\min}=0.02+11.83/L$, $R^2=0.9997$ — linearity in $1/L$ with four zero modes at every point.

A vortex in the AB-cloud at $\alpha=1/2$ behaves as a relativistic fermion in one dimension: the spectrum near zero is linear. The numerical check (Test 19) runs on lattices $L=12, 16, 24, 32, 48$: the minimal energy $E_{\min}$ (the gap at zero, with four zero modes) scales as $E_{\min}=0.0202+11.826/L$ with $R^2=0.9997$ — the linearity in $1/L$ is impeccable. Interpreting the constant: if the Dirac point sits at the Brillouin-zone center ($k_{\min}=2\pi/L$), the effective Fermi velocity is $v_F=b/2\pi\approx 1.88$ (in units of hopping over lattice constant b); if at the zone edge ($k_{\min}=\pi/L$), $v_F=b/\pi\approx 3.76$. The historical value ~ 0.125 belonged to a different lattice normalization; in the suite's current normalization we record precisely these two numbers, noting that choosing between them is a matter of identifying which band-contact point the zero-mode phase structure indicates. Test 31 (a torus variant with periodic boundaries) shows $R^2=0.03$ and zero gaps: on the torus the zero modes are topologically protected and never open at any L — itself a confirmation of the topological nature, but gap scaling requires open boundaries.

7.2 The Dirac Dip in the Density of States (Test 30)




Fig. 15. *The Dirac dip in the DOS at $\alpha=1/2$: a $20\times$ contrast in both directions (log scale).*

At $\alpha=1/2$ the Hofstadter spectrum has an analytic feature: the density of states near $E=0$ vanishes linearly (the Dirac point), whereas for $\alpha\neq 1/2$ the spectrum is generic. Test 30 measures the DOS in the central band $|E|<0.5$ on a 20×20 clean lattice: $\rho(\alpha=0.4)=0.200$, $\rho(\alpha=0.5)=0.010$, $\rho(\alpha=0.6)=0.200$ — a contrast of exactly $20\times$ in both directions. The Dirac dip is the fingerprint of the critical point separating Hall phases with $C_1=+1$ and $C_1=-1$; its observation in a lattice model (and analogues in cold atoms) is among the most robust evidences that genuine universal physics — not numerical tuning — is at work.

7.3 Hatano–Nelson: Leaving $\sigma=1/2$ (Test 32)




Fig. 16. Hatano–Nelson: $\sigma=0.5$ — real spectrum; $\sigma=0.7$ — the entire spectrum in the complex plane (skin effect).

The parameter σ in the generalized model sets the “fractional” part of the hopping phase; at $\sigma=1/2$ the model is Hermitian. Test 32 compares $\sigma=0.5$ and $\sigma=0.7$ on 16×16 : at $\sigma=0.5$, $\max |Im E|=0$ (the spectrum is strictly real); at $\sigma=0.7$ the entire 256-level spectrum is complex, $\max |Im E|=0.623$, and the Hatano–Nelson skin parameter is $g=(\sigma - 1/2) \ln |\gamma|=0.322$. The meaning:

leaving $\sigma=1/2$ moves the system from the unitary (GUE) class into the dissipative (uncorrelated) class; boundary localization of eigenvectors (the skin effect) destroys bulk statistics, and $\langle r \rangle$ on the non-Hermitian branch collapses to zero (all gaps “pair up” in the complex plane). The critical line $\sigma=1/2$ is singled out not only by Connes self-duality but as the only line where bulk random-matrix statistics exists at all (Figs. 13, 16). This is a second, independent argument for the “optimality of the critical line” — the statistical counterpart of the analytic self-duality.

7.4 Byte-Level Robustness (Test 37)

A practical reproducibility question: do the conclusions survive quantization? Test 37 verifies that $\langle r \rangle$ survives 256-level quantization of eigenvalues (change 0.0015 against tolerance 0.02) and that the generator’s byte stream has Shannon entropy 7.9991 bits out of a maximum of 8 (a CLT-based χ^2 criterion: $z=-0.234$, $|z|<3$). The test also records the global-vs-local classification difference: globally the spectrum of a particular realization

looks Poisson-like ($\langle r \rangle_{\rm raw} = 0.447$), but in the central window of 15 levels $\langle r \rangle_{\rm local} = 0.549$ — the bulk-vs-core distinction already noted in 5.4.

8. Block 6: Analytical Interpretations (Tests 21–23 and Extensions)

This block is the bridge from measurements to analytic structures. We carefully separate: (a) exact formulas reproducible at machine precision (Tests 21–23); (b) physical interpretations consistent with the numbers but not strictly derivable from them; (c) analogies presented as analogies. This three-level presentation is a matter of principle.

8.1 The Spinorial Phase γ^i : 90° Through the Imaginary Unit (Test 21)

Spinorial phase γ^* : argument $\approx \pi/2$




Fig. 17. The spinorial phase $\gamma^i = a_C + i b_C$: argument 89.87° against the reference 90° ; the dominance of b_C presses the phase to $\pi/2$.

Define $\delta_C = \pi/7 = 0.448799$ rad (see 8.2) and $b_2(K3) = 22$ (the second Betti number of the K3 surface). Test 21 verifies two real quantities: the **real braking term** $a_C(\gamma^{\dot{c}}) = \delta_C^5 / b_2(K3) = 8.276 \cdot 10^{-4}$ and the **imaginary Berry term** $b_C(\gamma^{\dot{c}}) = 1 - \cos(2\delta_C) = 2\sin^2\delta_C = 0.376510$; both reproduce at machine precision. The complex spinorial phase $\gamma^{\dot{c}} = a_C + i b_C$ has modulus $|\gamma^{\dot{c}}| = 0.3765$ and argument 89.874° — a deviation from $\pi/2$ of only 0.126° ($2.2 \cdot 10^{-3}$ rad). The structure is as follows: since $b_C \gg a_C$, the argument of $\gamma^{\dot{c}}$ is automatically pressed toward 90° — “the 90° transition is realized through the imaginary unit”: the phase, as a geometric object, is constructed so that its relativistic (spinorial) nature is expressed by purely imaginary dominance. The interpretation as a postulated Peccei–Quinn phase is at level (b): compatible with the numbers, not derived from the suite.

8.2 The AB Phase $\Phi_{AB} = \pi/7$ and the Fractional Charge (Test 22)

The exact formula: for a particle of charge $q/e = 1/14$ encircling a flux tube of one flux quantum, the Aharonov–Bohm phase is $\Phi_{AB} = 2\pi/14 = \pi/7 = 0.448799$ rad $\hookrightarrow 25.714^\circ$. Test 22 confirms the identity in exact rational arithmetic (discrepancy 0.00). The quantity $\delta_C = \pi/7$ is precisely Φ_{AB} ; that is, $q/e = 1/14$ is the postulate of fractional charge around which the entire analytic superstructure is built. Notably, $\pi/7$ connects to the Gaussian sum $b_7 = (-1 + i\sqrt{7})/2$ of the $\text{PSL}(2,7)$ character table (Section 8.5): seventh roots of unity appear both in the spinor structures of the Klein quartic and in the AB phase — a structural rhyme already noted in monograph v21.

8.3 The Complex Fractal Factor CF and Logarithmic Parabolization (Test 23)

For a critical solution with discrete scale ratio $\lambda_{\text{DSI}} = b_2(K3) = 22$, Test 23 computes the **complex fractal factor**

$$CF = \exp\left(\frac{2\pi i}{\lambda_{\text{DSI}}}\right) \cdot \exp\left(i \beta_{\text{imag}} \ln \lambda_{\text{DSI}}\right), \quad \beta_{\text{imag}} = 2 + \frac{1}{2\sin^2(\pi/7)} - 2\cos\frac{\pi}{7} = 2.854033,$$

yielding $CF = -0.9501 + 0.3119i$ with $|CF| = 1$ by construction. Note the logarithm inside the exponential — this is the “logarithmic parabolization”: modulating discrete scale invariance by a logarithmic phase. The test also carries an honest caveat: the hardcoded value in the late tex sources, $CF_{\text{hard}} = 0.9786 + 0.0390i$ (modulus 0.9794, whence $c_{AB} = 1 - |CF_{\text{hard}}| = 0.0206$), is an **empirical input** — an analogue of α_s in QCD — while the formula is the structural (topological) prediction; the amplitudes differ by exactly c_{AB} , and the test records both numbers with their statuses.

8.4 T-Violation as the Origin of GUE: The Black-Hole Analogy

The numbers of Blocks 3–5 assemble into Dyson’s analytic picture: $\text{GOE} \leftrightarrow \text{T-invariance}$, $\text{GUE} \leftrightarrow \text{its violation}$. The AB-cloud realizes both branches within one model: the Dirac gauge at integer charges — a real matrix, $\tau=0$, the GOE ceiling (Tests 15, 16); the smooth gauge or fractional charges — $\tau \approx 0.225$, GUE statistics (Tests 26, 33). A hypothesis proposed in discussions sharpens this line: **T-violation is the same reason the statistics of black holes (where time is warped and reversal is not a thermodynamic symmetry) is described by unitary ensembles**. We emphasize the status of this construction: it is not a consequence of the suite but an interpretive frame that the suite does not contradict — and that renders the prediction “any phase model with complex hoppings and no integrability will give $\langle r \rangle \rightarrow 0.6$ ” testable on independent material.

8.5 The Kerr Analogy, the Ergosphere, and $\text{PSL}(2,7) \subset \text{Monster}$ (Overview)

Two interpretations from v21 are retained in survey status. **The Kerr analogy**: the vortex phase field defines an “ergoregion” — a region where the phase velocity relative to the lattice changes sign, like the Kerr ergosphere; numerical verification requires dynamics rather than a static spectrum and remains a task. **The Monstrous bridge**: the Monster character χ_2 of dimension 196883, restricted to the subgroup $\text{PSL}(2,7)$ (embedded through the maximal subgroup $2^{3+6+12+18} \cdot (L_3(2) \times 3S_6)$), decomposes approximately as $1456 \cdot \chi_1 + 3334 \cdot \chi_3 + 3334 \cdot \chi_3' + 6784 \cdot \chi_6 + 8081 \cdot \chi_7 + 9770 \cdot \chi_8$ (sum of dimensions 196891, deviation 0.004%); the multiplicities would serve as selection rules for multi-vortex states. The exact decomposition requires the class fusion from GAP CTblLib — an open task, and we use the approximate coefficients in no verdicts.

9. The Large-Scale Vortex-Density Scan: From a Breeze to Ocean Currents (Test 34-Extension)

All chapters so far examined the AB-cloud in its reference configuration — a small number of vortices ($N_v = 2$) on a 96×96 lattice. This chapter reports a dedicated large-scale scan that asks the complementary question: given that the vortex configuration is an operator-valued gauge texture, which vortex parameters place the spectrum of the Hamiltonian closest to the statistics of the ζ zeros? The scan extends the direct-comparison protocol of Chapter 6 (the run-suite folder `test_34_direct_vs_zeta`; we refer to this extension as Test 34-ext) over all four adjustable dials of the construction — the vortex number N_v , the vortex charge q , the potential depth W and the flux α — together with the geometric choices (placement rule, boundary conditions), in three stages of increasing lattice size, and verifies the winners at the full 96×96 lattice under the composite acceptance criterion of Section 2.4. The

headline result is a broad, reproducible optimal window — $N_v \approx 8\text{--}16$ at $q \approx 0.7\text{--}1.0$ and the self-dual flux $\alpha = 0.5$ — in which the pooled KS distance to the unfolded ζ -zero spacings falls from $D = 0.1096$ (the suite’s reference value, WARN) to $D = 0.0318$: a factor of 3.4, within 0.010 of the finite-sample GUE floor of 0.0220. We offer three readings of this result: as a universal statement about vortex density in two-dimensional systems (Section 9.4), as an operator-theoretic working point of a tunable quantum system (Sections 9.2–9.3), and as numerical support for the statistical arm of the Hilbert–Pólya program (Section 9.5).

9.1 Why the Number of Vortices Matters — Across Physics

The vortex is the elementary carrier of phase texture in any two-dimensional medium, and the most economical classification of the states of such a medium is by vortex density — how many vortices per unit area the state carries. The same hierarchy of regimes recurs in every vortex-bearing system known to us, classical or quantum. A light breeze over a boundary layer is transported by a handful of coherent eddies: the vortex gas is dilute, every vortex is individually distinguishable, and the response of the flow depends on where exactly those few vortices happen to sit. Ocean-scale circulation — basin gyres and their mesoscale eddy fields — operates in the opposite limit: thousands of vortices per basin, strongly overlapping velocity and phase textures, transport governed by the collective field, and statistics that become reproducibly universal across oceans, models and decades precisely because no individual vortex matters any more. Between the two poles lie the familiar interpolated regimes: the Abrikosov lattice of a type-II superconductor passes from isolated flux tubes to overlapping cores as the external field raises the vortex density; a superfluid or a Bose–Einstein condensate crosses from potential flow to quantum turbulence when quantized vortices percolate through the sample. In all of these systems the response to vortex density is non-monotone: too few vortices cannot texture the medium at all, too many begin to re-organize into their own lattice or turbulent structures, and the working point lies in between, set by geometry and interactions.

The AB-cloud is a spectral incarnation of the same hierarchy. The vortices enter the Hofstadter Hamiltonian as an operator-valued gauge texture — Aharonov–Bohm phases on the bonds and on-site potential wells of depth proportional to qW — and the vortex number N_v is the density of that texture. The pure Hofstadter lattice ($N_v = 0$) is integrable-adjacent: every hopping phase is a rational multiple of the uniform flux α , the spectrum is far from GUE, and the KS distance to the ζ spacings is $D \approx 0.73$ (measured at $L = 24$; the cloud does not yet exist as a ζ -statistics object at all). A single vortex ($N_v = 1$) texturizes the operator only locally, $D \approx 0.19$; the reference pair ($N_v = 2$) reaches $D \approx 0.07\text{--}0.11$ depending on L and realization. The question of this chapter — at which vortex density the operator acquires the statistics of the ζ zeros — is therefore not rhetorical. The answer turns out

to be a window, not a point, and the window is broad, which is precisely what one wants from a construction that should be robust rather than fine-tuned.

9.2 The Operator System and the Three-Stage Scan Protocol

Every configuration of the scan is a finite-dimensional self-adjoint operator $H(N_v, q, W, \alpha; BC, \text{placement}; \text{seeds})$ — the Hofstadter Hamiltonian of Section 2.3 with the :monumental vortex phase model on an $L \times L$ lattice, with neutral vortex charges (first half $+q$, second half $-q$), vortex phases applied through the smooth gauge, and on-site wells $V = qW/(r^2/N + 1)$ — followed, unchanged, by the direct-comparison protocol of Chapter 6: five realizations per configuration (three in Stage A), spectra unfolded by the sliding-window rule, spacings compared with the first 50,000 ζ zeros by the two-sample Kolmogorov–Smirnov statistic D , complemented by $\langle r \rangle$, $R_2(s)$, the GUE/Poisson discriminator d_{GUE} vs d_{Pois} , and the composite acceptance criterion of Section 2.4 ($p > 0.01$ or $D < 0.10$; $\langle |\Delta R_2| \rangle < 0.10$; $d_{\text{GUE}} < d_{\text{Pois}}$; $\text{spread}(D) < 0.10$).

The scan proceeded in three stages. Stage A — a coarse sweep on $L = 24$ (576 sites): 171 configurations covering $N_v \in \{0, \dots, 20\}$, $q \in \{0.5, 0.7, 0.9, 1.0\}$, $W \in \{0, 0.5, 1.0, 1.5\}$, $\alpha \in \{0.3, 0.4, 0.5, 0.6\}$, three placement rules (protocol-random, regular grid, clustered) and two boundary conditions (open, torus). Stage B — refinement on $L = 32, 48, 64$: 46 configurations testing the finite-size stability of the Stage-A valley, the criticality of the flux α , and the per- L GUE floor (a GUE matrix of the same sample structure compared against the same ζ spacings — the statistical zero of the method). Stage C — verification on the full $L = 96$ lattice ($N = 9216$, 5528 spacings per realization, 27,640 pooled): five realizations each of the four leading candidates plus the reference configuration, under the full 34b protocol, with an independent $L = 96$ GUE control giving the finite-sample floor $D_{\text{floor}} = 0.0220$.

The computational engine is a faithful Python port of the suite’s Julia implementation, validated against the reference run before any scan data were accepted: per-realization unfolding counts match the reference log exactly (5528/5527 at $L = 96$), the ζ -side unfolding produces the same 49,999 spacings, the KS statistic agrees with the reference implementation to 10^{-12} , and cell flux on the torus is uniform to machine precision. The port’s own reference-configuration measurement ($D = 0.0928$, per-realization spread 0.072–0.111 over five realizations) reproduces the Julia suite’s level ($D = 0.1096 \rightarrow \text{WARN}$) within the run-to-run spread.

Determinism and memory were engineered explicitly: column-ordered matrices with in-place diagonalization, per-measurement subprocess isolation with `malloc_trim`, and a fixed seed protocol (base seed 96; vortex-placement RNG base + 3500, shared across realizations; disorder RNG base

+ 3500 + k per realization), so that every number in this chapter regenerates deterministically from the shipped scripts.

9.3 Results: the Optimal Vortex Window

Stage A located the valley. In N_v , the KS distance falls steeply from the pure lattice ($D = 0.73$) through $N_v = 1$ ($D = 0.194$) and $N_v = 2$ ($D = 0.073$) into a broad minimum at $N_v = 6$ – 16 ($D = 0.031$ – 0.065 , best at $N_v = 12$ on open boundaries: $D = 0.0393$ at $\alpha = 0.5$, $q = 1.0$, $W = 1.0$) and rises again beyond $N_v \approx 20$. In q the valley is 0.7 – 1.0 ; in W it is 0.5 – 1.0 ; in α the optimum at $L = 24$ sits at the self-dual value 0.5 , and Stage B shows that a deviation of α by ± 0.1 at $L = 32$ degrades D by roughly a factor of two. Geometry matters as much as the dials: clustered placements are catastrophically worse ($D = 0.11$ – 0.17 at $N_v = 2$ – 8), a regular grid is excellent at $N_v = 8$ ($D = 0.0385$), and the torus with protocol-random placement gives the best Stage-A surface overall ($N_v = 8$, torus: $D = 0.0319$). Figure 20 shows the $D(N_v, q)$ landscape in the two best W slices.




Fig. 20. The KS-distance landscape $D(N_v, q)$ on the $L = 24$ scan (Stage A; open boundaries, protocol placement, $\alpha = 0.5$): left — $W = 1.0$, right — $W = 0.5$. The valley $N_v \approx 8$ – 16 at $q \approx 0.7$ – 1.0 is visible in both slices; the pure lattice $N_v = 0$ sits at $D \approx 0.73$, far off the color scale.

Three facts in Table 1 carry the chapter. First, the reference configuration (row 1) sits at the very edge of the acceptance region: the port measures pooled $D = 0.0928$ with a per-realization spread of 0.072 – 0.111 that straddles the $D = 0.10$ boundary, and the Julia suite’s own reference run reports $0.1096 \rightarrow \text{WARN}$. Second, all three optimized candidates pass the composite criterion with margin, and the best of them ($N_v = 16$, $q = 1.0$, $\alpha = 0.5$, torus) reaches $D = 0.0318$ — a factor of 3.4 below the suite’s reference value. Third, the best candidate has essentially exhausted the statistical room available to any GUE-convergent sample of this size: the floor itself is $D = 0.0220$, so the cloud at 0.0318 closes $\approx 89\%$ of the reference-to-floor gap. The residual 0.010 is not necessarily a defect of the construction: the floor itself fails the p-arm of the criterion ($p = 3.0 \cdot 10^{-4}$), a finite-ensemble effect of exactly the kind discussed in Section 3.3 — and Figure 21 shows

that the optimized configurations hug the per-L floor all the way up the size ladder, while the reference configuration remains roughly three times above it.

Table 1. Stage-C verification on the full 96×96 lattice (five realizations per configuration, 27,640 pooled spacings; ζ side: 49,999 unfolded spacings of the first 50,000 zeros; composite criterion of Section 2.4). D is the KS distance pooled over realizations; Spread is the min-max range of per-realization D . The GUE control row is the finite-sample floor: the same KS comparison applied to a GUE matrix ensemble of the same sample structure.

Configuration	D	Spread	$\langle \Delta R^2 \rangle$	dGUE	Verdict
$N_v = 2, q = 1.0, \alpha = 0.5$, open (reference)	0.0928	0.072–0.111	0.0549	0.0454	edge
$N_v = 8, q = 1.0, \alpha = 0.6$, open, regular grid	0.0358	0.033–0.039	0.0487	0.0309	PASS
$N_v = 16, q = 0.7, \alpha = 0.5$, open	0.0387	0.037–0.041	0.0336	0.0231	PASS
$N_v = 16, q = 1.0, \alpha = 0.5$, torus (best)	0.0318	0.028–0.038	0.0357	0.0218	PASS
GUE control (same sample structure)	0.0220	—	0.0272	—	floor

The mechanism of the window is phase-texture coverage. In the dilute regime ($N_v \lesssim 4$) large coherent patches of the lattice retain nearly real, nearly integrable hopping structure; $\langle r \rangle \approx 0.51$ – 0.58 and $D > 0.07$. In the working window ($N_v \approx 8$ – 16) the vortex phase fields overlap and jointly complexify every bond of the operator — precisely the regime in which the BGS prescription (a Hermitian operator with no antiunitary symmetry and complex phases) is actually realized — and D reaches the floor. In the dense regime ($N_v \gtrsim 20$) the on-site wells begin to overlap into their own potential landscape and the vortex gas re-organizes, degrading the uniformity of the texture again — the same self-obstruction that an over-dense vortex gas displays in fluids. The charge dial confirms the reading: $q = 1$ completes the flux quantization $2\pi q = 2\pi$ and maximizes the phase incommensurability

across the texture, while fractional charges leave partially commensurate phases (and, per Section 4.2, remain physically visible — Byers–Yang). The flux dial echoes the monograph’s earlier theme: $\alpha = 0.5$ is the self-dual critical line (Sections 4.3 and 8.1), and it is strictly optimal here as well.




Fig. 21. The size ladder $L = 24 \rightarrow 64$ for the six leading Stage-B configurations (pooled D per lattice size; solid curves) against the per- L GUE floor (dashed). The optimized window stays at or near the floor as the lattice grows; the dilute configurations plateau roughly three times above it.

9.4 From a Breeze to Ocean Currents: the Physical Reading

The scan gives the vortex-density hierarchy of Section 9.1 quantitative content in the spectral setting. The dilute regime of the cloud — $N_v = 0-2$ — is its breeze: a few identifiable vortical structures, a phase texture that is still nearly regular, a response that depends strongly on the individual configuration, and no universality — $D \approx 0.09-0.11$ with a wide per-realization spread. The optimal window $N_v \approx 8-16$ is the cloud’s gyre regime: the vortex fields overlap into a collective texture, individual placement stops mattering (the spread of D over realizations shrinks accordingly), and the statistics become universal in the strict sense available to a finite quantum system — within 0.010 of the GUE floor measured against the ζ zeros. The dense regime reproduces the fluid’s self-obstruction from the opposite side. The correspondence is a structural one — the AB-cloud’s vortices are gauge vortices, phase singularities of a quantum wavefunction, not fluid eddies — but the control parameter and the shape of the response are the same: a non-monotone dependence of the order measure on vortex density, with a broad working plateau in between.

Stated as an operator-theoretic fact: the scan maps a one-parameter family of quantum systems — the Hamiltonian family $H(N_v, \dots)$ — onto a scalar distance-to- ζ , and the minimum of that function is both broad and shallow, the signature of a robust working point rather than a tuned one. In fluid language: the construction works not because someone blew exactly the right gust, but because somewhere between a breeze and an ocean current there is a wide regime in which vortices collectively do their transport job; the cloud’s job is phase transport toward GUE universality, and $N_v \approx 8\text{--}16$ is where that transport saturates. We emphasize the honest boundary of the analogy: what transfers is the density-controlled crossover from individual to collective vortex behavior; no literal hydrodynamics is claimed, and none is needed.

9.5 The Hilbert–Pólya Perspective

The Hilbert–Pólya program seeks a self-adjoint operator whose spectrum is the set of ζ zeros. The AB-cloud family is not that operator — its eigenvalues are not the zeros, and nothing in this chapter claims otherwise. What the scan establishes is a necessary condition any Hilbert–Pólya candidate must satisfy, verified here in the strongest available statistical form: a Hermitian operator of a simple, physically transparent gauge type — a lattice Hamiltonian whose phases are generated by quantized gauge vortices — can reproduce the spacing statistics of the ζ zeros to within the finite-sample floor ($D = 0.0318$ against 0.0220 ; $\approx 89\%$ of the reference-to-floor gap closed). This is numerical support for the statistical, Montgomery–Odlyzko arm of the program: the ζ statistics are not merely GUE-shaped in outline; they are approachable, with quantitative closeness in the KS metric, by operators of vortex-gauge construction.

The scan also sharpens what “reference-grade” means for the construction as a whole. Every dial of the vortex system now has a measured response curve: N_v (non-monotone, optimum $8\text{--}16$), q (optimum at full charge, valley $0.7\text{--}1.0$), W (plateau $0.5\text{--}1.0$), α (strictly critical at the self-dual 0.5), placement (uniform or regular good, clustered destructive), boundary conditions (torus best, open acceptable). Every response carries its seeds, its spread and its floor. In the taxonomy of Chapter 11 — proven, reliably measured, open — the assignment is: proven, that the port reproduces the suite’s reference quantities exactly and the 34b protocol is implemented faithfully; reliably measured, the D values of Table 1 with their spreads and the GUE floor; open, whether the window persists at $L = 128$ and against the 500,000-zero Odlyzko list already present in the repository, and whether an analytic reason for the width of the window — rather than the numeric valley — can be given.

The relation to the Hilbert–Pólya conjecture must be stated with the discipline we apply everywhere else in this monograph: the result is evidence, not proof. No arithmetic information about individual zeros enters

the comparison — only unfolded spacings — so nothing here touches the additivity, multiplicativity or simplicity of the zeros, and the Riemann hypothesis remains exactly as open as before. What the result does establish, and what we submit as the construction’s reference-grade credential, is that the statistical target of the Hilbert–Pólya program is reachable by a verified, seed-deterministic, fully reproducible quantum operator system whose every parameter is documented.

9.6 Reproducibility

All scan artifacts ship with the repository in the folder `t34_vortex_optimization/`: the three sweep tables (`sweep_A_L24.csv`, `sweep_B_refine.csv`, `sweep_C_L96.csv` — one row per configuration with the full metric set), the GUE control (`gue_control_L96.json`), the $R_2(s)$ curves of the four Stage-C configurations (npz), the seven report figures, and the scripts (`t34_engine.py` — the validated port; `t34_sweep.py` — the staged scan driver; `gue96.py` — the $L = 96$ GUE control; `t34_charts.py` — the figures; `test_engine.py` — the port-validation suite). A single Stage-C configuration (five realizations at $L = 96$) takes ≈ 13 minutes on a laptop-class core (≈ 155 s per realization, peak memory 1.4 GB); the whole Stage A takes under two minutes. The PDF report in the same folder reproduces this chapter with the complete tables, and Table 1 here is regenerated directly from `sweep_C_L96.csv`.

The recommended upgrade for the Julia suite itself follows from row 4 of Table 1: switching the reference run to $N_v = 16$ on the torus (the flag `--ab-nv-secondary 16` with periodic boundaries, or the equivalent configuration keys) moves the suite’s headline statistic from the WARN band into the deep PASS band without touching any other part of the protocol. We deliberately leave the suite’s default untouched: the reference configuration $N_v = 2$ remains valuable as the conservative, dilute-regime benchmark against which the optimization is measured.

10. Synthesis: ζ Zeros as Codes of Admissible States




Fig. 18. The phase texture of a vortex pair $q=\pm 1$ over the Landau gauge: the phase “cloud” that generates the observed statistics.




Fig. 19. *The Hofstadter butterfly $E(\alpha)$; the vertical marks the self-dual point $\alpha=1/2$ where the key results of the monograph concentrate.*

Assemble the chain. (1) The spectrum of the AB-cloud at $\alpha=1/2$ in the smooth gauge has GUE statistics converging with size: $\langle r \rangle : 0.548 \rightarrow 0.595$ for $L : 10 \rightarrow 50$ (Tests 33, 34). (2) The statistics of the ζ zeros is GUE plus computable finite-sample corrections (Tests 4–14); at reachable heights “pure GUE” is rejected, “finite-size GUE” is not (Test 12). (3) The direct comparison shows the qualitative structure of R_2 matches (Test 35) — the Montgomery correlation hole is reproduced by the vortex cloud. (4) The critical line $\sigma=1/2$ is simultaneously the point of Connes self-duality, the Dirac dip, and Hermiticity (Tests 17, 30, 32) — and historically the point where the cloud’s KS distance to the zeros is minimal. (5) Vortices behave as relativistic particles (the Dirac cone, Test 19), and fractional charge acts as physical flux (Byers–Yang, Test 25). The final narrative — “ ζ zeros as codes of admissible energy states of a phase resonator” — thereby attains the status of a **consistent, quantitatively verified program**, not a proven theorem: we know which links are machine-exact (fluxes, identities), which

are statistically robust (scaling, bootstrap), and which remain open (the $T \rightarrow \infty$ asymptotics, pointwise agreement of R_2).

11. Verification Methodology: What a “PASS” Actually Guarantees

Suite version 19 is the first release in which the report engine is architecturally isolated from the computational core. The history of this decision is instructive: in v18, a bug in the PDF writer (the line list was typed as `String[]` while tuples were pushed into it) raised a `MethodError` after the completion of **every** test; the exception was caught by the outer handler, verdicts were downgraded to `WARN`, and the second pass of the two-pass mode never ran at all — the “quick check” degraded to single-pass while outwardly appearing to work. Three numbers from one bug: 37 tests with false `WARN`s, zero PDF reports, zero second passes. The fix (FIX-R1... R5 in Appendix A) established: (1) the PDF writer is correct; (2) the failure of any report format cannot affect a verdict; (3) Test 28 (variable scope of the two-pass loop) and Test 18 (torus gauge compatibility at $\alpha=1/3$ on non-multiples of six) are fixed; (4) the quick check is truncated to 5000 zeros and genuinely performs both passes $16 \times 16 \rightarrow 32 \times 32$. The lesson: **a suite’s conclusions are no stronger than its weakest infrastructure**, and an “almost working” report pipeline is more dangerous than a broken one because it fabricates an appearance of verification.

Practical guarantees of the current version: every test is a separate directory with full logs; every number of the monograph traces to a test ID; a verdict is reproduced by a command of the form `julia ab_cloud_v19.jl --test N --no-two-pass`; two-pass mode is on by default for the full run and for `--quick`; the RNG uses fixed seeds everywhere randomness is needed (bit-for-bit reproducibility).

12. Conclusions

1. **Construction exactness** (machine level): Hermiticity 0; Dirac-string fluxes to 10^{-14} ; Byers–Yang for integer charges $3.5 \cdot 10^{-15}$; Connes self-duality — 4 zero modes and $E \rightarrow -E$ at defect 10^{-15} ; chirality 0.0 at $W=0$. The construction is mathematically correct — proven, not estimated.
2. **GUE universality** (statistical level): $\langle r \rangle = 0.5848 \pm 0.0260$ against GUE 0.5992 (deviation -2.4%); scaling $L=10 \rightarrow 50$ is monotone and plateaus at 0.595 ± 0.008 ; the spectrum is stiffer than Poisson at all scales L and closer to GUE in 9/9 points. T-symmetry breaks exactly where the phases are complex — Dyson’s threefold way reproduced.
3. **Comparison with the ζ zeros** (honest level): the qualitative structure of $R_2(s)$ matches (hole, plateau; closer to GUE than Poisson), pointwise agreement is absent and should not be expected at 304

gaps; Berry's finite-sample corrections quantitatively explain all observed deviations from asymptotic GUE.

4. **Topology and dynamics:** $C_1=2$ (non-trivial IQHE), the Dirac cone $R^2=0.9997$, the DOS dip $20\times$, the skin effect at $\sigma\neq 1/2$; the critical line $\sigma=1/2$ is triply singled out (self-duality + Hermiticity + dip).
5. **Analytical interpretations:** the formulas for γ^i , $\Phi_{AB}=\pi/7$, and CF reproduce at machine precision and form a coherent $\pi/7$ -structure; their physical interpretation (Peccei-Quinn, Moonshine) remains hypothetical — the suite does not contradict it, but does not prove it.
6. **Infrastructure of trust:** v19 repairs every discovered defect of the report pipeline; every conclusion traces to logs; the two verdict systems (machine and calibrated) are explicitly separated.

Finally, the construction itself receives its quantitative verdict from the large-scale scan of Chapter 9: with the vortex dial set to its verified working window ($N_v \approx 8-16$, $q \approx 1.0$, $\alpha = 0.5$, torus), the AB-cloud's spacing statistics stand at $D = 0.0318$ against the ζ zeros — $3.4\times$ closer than the suite's reference configuration and within 0.010 of the finite-sample GUE floor. A construction whose optimal parameters are sought, found, verified and documented to this level of reproducibility is what we mean by a reference-grade quantum operator system.

13. Open Questions

1. **The asymptotics of $b(N)$:** power law or logarithm? The alternative fit is better ($R^2=0.9999$), but the range of N is insufficient; a run to $N \sim 10^6$ is needed (the 2M Odlyzko dataset is already supported by the loader).
2. **Pointwise agreement of R_2 :** the AB-cloud ensemble must grow to $\geq 10^4$ gaps (larger lattices $\times$ multiple realizations $\times$ vortex shifts) — only then does the two-sample KS against ζ -5000 become informative.
3. **The form factor $K(t)$:** same goal; the PASS thresholds ($RMS < 0.30$, $corr > 0.50$) are, by our estimate, reachable at ≥ 3000 gaps.
4. **The 3D interface:** vortex lines in a Hofstadter stack (the 3D-lab module) — does the GUE plateau persist along the third dimension; preliminary data exist but no 37-test verification.
5. **The exact decomposition of 196883 over PSL(2,7):** requires the class fusion via GAP CTblLib; until then the coefficients remain approximate.
6. **Vortex dynamics and the Kerr analogy:** a static spectrum cannot test the ergosphere interpretation; a Lagrangian simulation of point vortices with AB phases is needed.

The vortex-density scan of Chapter 9 opens a question of its own. The optimal window $N_v \approx 8-16$ reaches $D = 0.0318$ but does not cross the finite-

sample GUE floor of 0.0220: is the residual gap a property of the cloud's sample structure, or a genuine statistical signature of the ζ zeros beyond GUE (cf. the rigidity excess of Section 3.4)? The direct next measurements are already possible with the repository's assets — the Odlyzko 500,000-zero list for the ζ side, and $N_v \in \{16, 24, 32\}$ at $L = 128$ for the cloud side.

Appendix A. The Code: Architecture and Key Modules

The full code of the verification suite (file `ab_cloud_v19.jl`, Julia 1.10, zero external dependencies — including its own implementations of a PNG writer with pHYs, vector PDF, GIF89a with LZW, DOCX as stored-ZIP, and all of the statistics) ships with the monograph in the code/ directory; run without arguments it opens an interactive menu (key `f` — the quick check). Below we list the architectural modules and quote the key listings.

Modules of the file: (1) the Config structure and bilingual messages; (2) data loaders — 50,000 embedded Odlyzko zeros + external files up to 2M with graceful fallback; (3) the statistical core (unfolding, KS, χ^2 , AD, bootstrap, Σ^2 , Δ_3 , R_2 , $K(t)$); (4) the AB-cloud builder `build_ab_cloud_hamiltonian` with the phase models (`:dirac`, `:monumental`) and the flux check `plaquette_flux`; (5) the 37 tests `test*`; (6) the two-pass runner `run_test_two_pass`; (7) the v16/v19 report engine (`rep_*`, the `md/html/pdf/docx/svg/png/gif` writers, `index.html`); (8) the physics lab (22 experiment configurations) and the 3D lab (30 tests); (9) the menus and CLI.

The key function — the Hamiltonian builder (abridged; the full text is in the shipped file):

```
#  $\varphi_{ij} = 2\pi\alpha \cdot j + \sum_k q_k [\arg(r_{i-r_k}) - \arg(r_{j-r_k})] + \varphi^{\text{string}_{ij}}$ 
# bond (i→j) receives  $H_{ij} = -\exp(i\varphi_{ij})$ ; a plaquette with vortex k carries flux  $2\pi \cdot q_k$ 
function build_ab_cloud_hamiltonian(cfg::ABCloudConfig)
    N = cfg.Nx * cfg.Ny
    H = zeros{ComplexF64, N, N}
    for iy in 0:(cfg.Ny-1), ix in 0:(cfg.Nx-1)
        i = iy * cfg.Nx + ix
        for (dx, dy) in ((1,0),(0,1)) # nearest neighbors, both orientations
            jx, jy = (ix+dx) % cfg.Nx, (iy+dy) % cfg.Ny
            j = jy * cfg.Nx + jx
             $\varphi = 2\pi * \text{cfg.alpha} * jy$  # Landau gauge (column phase)
            for v in cfg.vortices # vortex phases (model-dependent)
                 $\varphi += \text{vortex\_phase}(v, ix, iy, jx, jy, \text{cfg.model})$ 
            end
            H[i+1, j+1] -= exp(im *  $\varphi$ )
            H[j+1, i+1] -= exp(-im *  $\varphi$ ) # Hermitian conjugate
        end
        H[i+1, i+1] += disorder_epsilon(ix, iy, cfg) #  $\varepsilon \sim U(-W/2, W/2)$ , fixed seed
    end
end
```

```
    return H
end
```

The plaquette flux check (Dirac gauge): the sum of the four bond phases of a plaquette containing the vortex gives $2\pi q \pmod{2\pi}$, an empty one — 0; computed by `plaquette_flux(H, Nx, Ny, ix, iy)` and verified in Tests 15/24/26.

Verdicts: `log_result` writes the verdict line and the full journal; `run_test_two_pass` runs the test on the primary and the enlarged lattice, restoring the configuration in `finally` (pass isolation); `rep_finish!` records the verdict BEFORE writing the reports, and a writer failure (FIX-R2) cannot alter it.

The v19 fixes (full comments in the code): FIX-R1 — typing of the intermediate line list in the PDF writer (eliminates the mass `MethodError` that nullified verdicts and the second pass); FIX-R2 — `try/catch` around `_write_test_report!`; FIX-R3/3b — `last_R2/berry_ref` in the function scope of Test 28; FIX-R4 — lattice snap to a multiple of 6 for $\alpha \in \{1/2, 1/3\}$ in Test 18; FIX-R5/5b — quick check: ζ truncation to 5000, GIF disabled, the `--quick` CLI flag.

Appendix B. Reproduction Guide

the full suite (two-pass, many minutes):

```
julia ab_cloud_v19.jl --test all
```

a single test:

```
julia ab_cloud_v19.jl --test 33 --no-two-pass
```

the quick check $16 \times 16 \rightarrow 32 \times 32$ ($\zeta \leq 5000$, minutes):

```
julia ab_cloud_v19.jl --quick
```

language:

```
julia ab_cloud_v19.jl --lang en --quick
```

reports: results/run_<date>/test_NN_<slug>/{report.md,html,pdf,docx, plots, logs/}

The numbers of the monograph correspond to the full two-pass run **v19 run_20260902_134759** of 2026-09-02 (37 tests, two-pass $72 \times 72 \rightarrow 96 \times 96$ with the **HARDCORE** second-pass audit, a pass-3 series for tests 19/29/31; Julia 1.12.0; 50,000 Odlyzko zeros). First-pass summary: **32 PASS / 5 WARN** (tests 4, 5, 6, 11, 12, 28, 35 — statistical WARN verdicts of the “Wigner-surmise floor”, calibrated by two-sample KS lines against the exact GUE reference; suite semantics unchanged); **all HARDCORE second-pass runs are PASS** (77 verdict lines in total). The full artifact set — per-test reports (report.md/pdf/docx/html), computation logs (logs/), the final `FINAL_REPORT` and `index.html` — ships in the repository: `results/run_20260902_134759/`. In addition, all 64 spinor structures of the Klein quartic were verified by the reference implementation `verification/spinor64` (Appendix D). The early single-pass v18 baseline (`results/verification_run_v18_37tests_2026-08-28.txt`) is kept for history.

Appendix C. Figure Index

All figures are 600 dpi PNG with captions in the language of the edition; the files live in monographs/<lang>/figures/. Fig. 1 — $b(N)$ convergence (two fits); Fig. 2 — ζ spacing histogram vs. Wigner surmises; Fig. 3 — the two decay fits in log coordinates; Fig. 4 — bootstrap distribution of the slope with the 95% CI; Fig. 5 — KS $D(T_{min})$ with the critical line; Fig. 6 — $\Sigma^2(L)$; Fig. 7 — $\Delta_3(L)$; Fig. 8 — $R_2(s)$: AB-cloud vs. ζ , GUE, Poisson; Fig. 9 — the form factor $K(t)$; Fig. 10 — scaling $\langle r \rangle(L)$ for $q=0.3$ and $q=1$; Fig. 11 — the $\langle r \rangle$ bootstrap (Test 33); Fig. 12 — the Dirac cone $E_{min}(1/L)$; Fig. 13 — the Dirac dip in the DOS; Fig. 14 — Byers-Yang (integer vs. fractional charge); Fig. 15 — the Berry cutoff $R_2(0;T)$; Fig. 16 — Hatano-Nelson, the ellipse of the non-Hermitian spectrum; Fig. 17 — the spinorial phase γ^6 ; Fig. 18 — the phase texture of a vortex pair; Fig. 19 — the Hofstadter butterfly with the point $\alpha=1/2$ marked.

Appendix D. Verification of all 64 spinor structures of the Klein quartic (v22.1, spinor64)

The reference implementation verification/spinor64 (Python/NumPy, reproduction: python3 verification/spinor64/run_spinor64.py) performs two independent experiments over all 64 spinor structures $\varepsilon \in F_2^6$ of the Klein quartic and corrects the v21 monograph claim about the “uniqueness” of $\text{idx}=38$.

E1 — exact symmetry (the Klein graph {3,7}). The Klein quartic is discretised by its regular map {3,7} (56 vertices, 84 edges, 24 heptagons, $\text{Aut} = \text{PSL}(2,7)$ of order 168). The 64 spin structures are realised as edge signings with odd face parity (the Kasteleyn/Cimasoni-Reshetikhin model, canonical gauge: all spanning-tree edges positive). The computed splitting of the $\text{PSL}(2,7)$ action gives orbits **28 / 21 / 7 / 7 / 1**; the 28-element orbit — the odd structures ($\text{Arf}=1$) — is a single orbit, confirming the transitivity of $\text{PSL}(2,7)$ on the 28 bitangents (the Riemann-Klein theorem). The signed Dirac operators within an orbit are permutation-conjugate: the maximum pairwise spectral distance over all 64 structures is **$8.9 \cdot 10^{-15}$** ; gauge invariance $7.1 \cdot 10^{-15}$; zero modes of the discrete operator: 2 (odd orbit) / 3 (even orbits) / 7 (trivial class). Consequence: no spinor structure can carry unique statistics — the “uniqueness of $\text{idx}=38$ ” in v21 was a discretization artifact breaking the $\text{PSL}(2,7)$ symmetry. By the v21 monograph’s own formula $\text{Arf}(\varepsilon) = \varepsilon_1 \varepsilon_2 + \varepsilon_3 \varepsilon_4 + \varepsilon_5 \varepsilon_6$ the vector $\varepsilon(38) = (0, 1, 1, 0, 0, 1)$ gives $\text{Arf} = 0$ (not 1 as previously claimed).

E2 — statistics (AB-cloud). Hofstadter torus $L=44$, $\alpha=1/2$, $N_v=54$ vortices $q=+1$ (density-scaled anchor $N_v=25$ at 30×30), $W=0$, the :monumental gauge

(atan smooth, vertical bonds, factor 0.5), spin structure as boundary twists $\varphi_x = \pi(\varepsilon_1 + \varepsilon_3 + \varepsilon_5)$, $\varphi_y = \pi(\varepsilon_2 + \varepsilon_4 + \varepsilon_6)$; averaging over 5 vortex configurations, bulk window 0.6. Reference: a Monte-Carlo GUE ensemble of 100 matrices 1936×1936 (the Test-16 methodology): median $\langle r \rangle = 0.6013$, 95% CI $[0.5847; 0.6140]$. Result: **all 64 structures lie inside the CI**: $\langle r \rangle = 0.5984 \pm 0.0035$ (spread 0.5935...0.6027, min MC p = 0.36). The GUE statistics is identical for all structures — the source of GUE, as concluded in Section 4.1, is the AB-cloud dynamics, not the choice of spinor structure.

Ports of Test 38 (frozen data + a self-implemented Jacobi algorithm, no LAPACK) ship in 10 languages:
[verification/{cpp,java,rust,go,fortran,haskell,r,matlab,julia,javascript}/spinor38/](#).